

Fermilab Keras Workshop

Stefan Wunsch
stefan.wunsch@cern.ch

November 30, 2017

What is this talk about?

- ▶ Modern implementation, description and application of neural networks
- ▶ Currently favoured approach:
 - ▶ **Keras** used for **high-level description** of neural networks models
 - ▶ **High-performance implementations** provided by backends, e.g., Theano or **TensorFlow** libraries

Being able to go from idea to result with the least possible delay is key to doing good research.

theano



Outline

The workshop has these parts:

1. Brief introduction to **neural networks**
 2. Brief introduction to **computational graphs** with TensorFlow
 3. Introduction to **Keras**
 4. **Useful tools** in combination with Keras, e.g., TMVA Keras interface
- ▶ In parts 3 and 4, you have to possibility to follow along with the examples on your laptop.

Assumptions of the tutorial:

- ▶ You are not a neural network expert, but you know roughly how they work.
- ▶ You haven't used Keras before.
- ▶ You want to know why Keras is so popular and how you can use it!

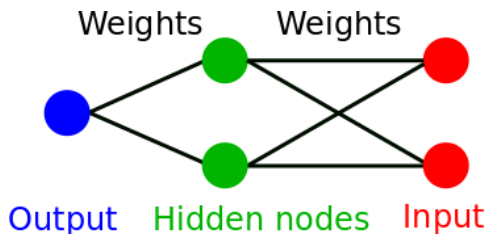
You can download the slides and code examples from GitHub:

`git clone`

`https://github.com/stwunsch/fermilab_keras_workshop`

Brief Introduction to Neural Networks

A Simple Neural Network



Neural Network: $f(x)$

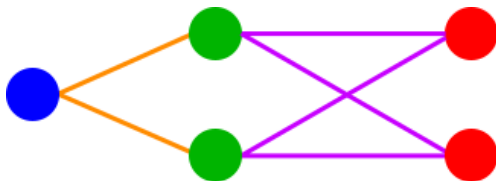
- ▶ **Important:** A neural network is only a mathematical function. No magic involved!
- ▶ **Training:** Finding the best function for a given task, e.g., separation of signal and background.

Mathematical Representation

- **Why do we need to know this?**

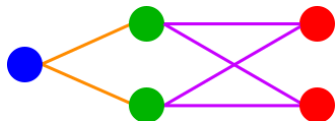
- Keras backends TensorFlow and Theano implement these mathematical operations explicitly.

- Basic knowledge to understand Keras' high-level layers



$$f_{\text{NN}} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$

Mathematical Representation (2)



$$f_{\text{NN}} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$

$$\text{Input : } x = \begin{bmatrix} x_{1,1} \\ x_{2,1} \end{bmatrix}$$

$$\text{Weight : } W_1 = \begin{bmatrix} W_{1,1} & W_{1,2} \\ W_{2,1} & W_{2,2} \end{bmatrix}$$

$$\text{Bias : } b_1 = \begin{bmatrix} b_{1,1} \\ b_{2,1} \end{bmatrix}$$

Activation : $\sigma(x) = \tanh(x)$ (as example)

Activation is applied elementwise!


The “simple” neural network written as full equation:

$$f_{\text{NN}} = \sigma_2 \left([b_{1,1}^2] + [W_{1,1}^2 \quad W_{1,2}^2] \sigma_1 \left(\begin{bmatrix} b_{1,1}^1 \\ b_{2,1}^1 \end{bmatrix} + \begin{bmatrix} W_{1,1}^1 & W_{1,2}^1 \\ W_{2,1}^1 & W_{2,2}^1 \end{bmatrix} \begin{bmatrix} x_{1,1} \\ x_{2,1} \end{bmatrix} \right) \right)$$

► How many parameters can be altered during training?

→ $1+2+2+4=9$ parameters

Training (Short Reminder)



$$f_{NN} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$

Training:

1. **Forward-pass** of a **batch** of N inputs x_i calculating the outputs $f_{NN,i}$
2. **Comparison of outputs** $f_{NN,i}$ with true value $f_{Target,i}$ using the **loss function** as metric
3. **Adaption of free parameters** to improve the outcome in the next forward-pass using the gradient from the **back-propagation** algorithm in combination with an **optimizer algorithm**

Common loss functions:

- ▶ Mean squared error: $\frac{1}{N} \sum_{i=1}^N (f_{NN,i} - f_{Target,i})^2$
- ▶ Cross-entropy: $-\sum_{i=1}^N f_{Target,i} \log(f_{NN,i})$

Deep Learning Textbook

Free textbook written by Ian Goodfellow, Yoshua Bengio and Aaron Courville:

<http://www.deeplearningbook.org/>

- ▶ Written by leading scientists in the field of machine learning
- ▶ **Everything you need to know** about modern machine learning and deep learning in particular.

▶ Part I: Applied Math and Machine Learning

Basics

- ▶ 2 Linear Algebra
- ▶ 3 Probability and Information Theory
- ▶ 4 Numerical Computation
- ▶ 5 Machine Learning Basics

▶ II: Modern Practical Deep Networks

- ▶ 6 Deep Feedforward Networks
- ▶ 7 Regularization for Deep Learning
- ▶ 8 Optimization for Training Deep Models
- ▶ 9 Convolutional Networks
- ▶ 10 Sequence Modeling: Recurrent and Recursive Nets
- ▶ 11 Practical Methodology
- ▶ 12 Applications

▶ III: Deep Learning Research

- ▶ 13 Linear Factor Models
- ▶ 14 Autoencoders
- ▶ 15 Representation Learning
- ▶ 16 Structured Probabilistic Models for Deep Learning
- ▶ 17 Monte Carlo Methods
- ▶ 18 Confronting the Partition Function
- ▶ 19 Approximate Inference
- ▶ 20 Deep Generative Models

Brief Introduction to Computational Graphs With TensorFlow

Motivation

- ▶ Keras wraps and simplifies usage of libraries, which are optimized on efficient computations, e.g., TensorFlow.
- ▶ How do modern numerical computation libraries such as Theano and TensorFlow work?

Theano? TensorFlow?

- ▶ **Libraries for large-scale numerical computations**
- ▶ TensorFlow is growing much faster and gains more support (Google does it!).

 [Theano](#) / [Theano](#)

 Watch ▾

512

★ Star

5,893

 Fork

2,031

Theano** is a Python library that allows you to **define, optimize, and evaluate mathematical expressions involving multi-dimensional arrays efficiently.

 [tensorflow](#) / [tensorflow](#)

 Watch ▾

4,641

★ Unstar

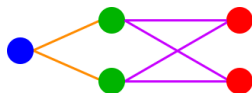
50,822

 Fork

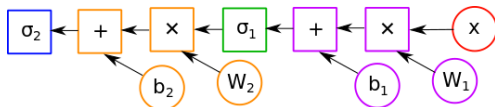
23,745

***TensorFlow** is an open source software library for **numerical computation using data flow graphs**. Nodes in the graph represent mathematical operations, while the graph edges represent the multidimensional data arrays (tensors) communicated between them.*

Computational Graphs



$$f_{NN} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$



Example neural network




According computational graph

- ▶ TensorFlow implements all needed **mathematical operations for multi-threaded CPU and multi GPU** environments.
- ▶ Computation of neural networks using data flow graphs is a perfect match!

***TensorFlow** is an open source software library for numerical computation using data flow graphs. **Nodes** in the graph **represent mathematical operations**, while the **graph edges represent the multidimensional data arrays (tensors)** communicated between them.*

TensorFlow Implementation of the Example Neural Network



`fermilab_keras_workshop/tensorflow/xor.py:`

```
w1 = tensorflow.get_variable("W1", initializer=np.array([[1.0, 1.0],
                                                         [1.0, 1.0]]))
b1 = tensorflow.get_variable("b1", initializer=np.array([0.0, -1.0]))
w2 = tensorflow.get_variable("W2", initializer=np.array([[1.0], [-2.0]]))
b2 = tensorflow.get_variable("b2", initializer=np.array([0.0]))

x = tensorflow.placeholder(tensorflow.float64)
hidden_layer = tensorflow.nn.relu(b1 + tensorflow.matmul(x, w1))
y = tensorflow.identity(b2 + tensorflow.matmul(hidden_layer, w2))

with tensorflow.Session() as sess:
    sess.run(tensorflow.global_variables_initializer())
    x_in = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
    y_out = sess.run(y, feed_dict={x:x_in})
```

→ **Already quite complicated for such a simple model!**

TensorFlow Implementation of the Example Neural Network (2)

- ▶ Plain TensorFlow implements only the mathematical operations.
- ▶ Combination of these operations to a neural network model is up to you.
- ▶ Already quite complicated for a simple neural network model without definition of loss function, training procedure, ...

- ▶ **Solution 1:** Write your own framework to simplify TensorFlow applications
- ▶ **Solution 2:** Use wrapper such as Keras with predefined layers, loss functions, ...

Introduction to Keras

What is Keras?


- ▶ Most popular tool to train and apply (deep) neural networks
- ▶ **Python wrapper around multiple numerical computation libraries**, e.g., TensorFlow
- ▶ Hides most of the low-level operations that you don't want to care about.
- ▶ **Sacrificing little functionality** for much easier user interface
- ▶ **Backends:** TensorFlow, Theano
- ▶ **NEW:** Microsoft Cognitive Toolkit (CNTK) added as backend


theano





Why Keras and not one of the other wrappers?


- ▶ There are lot of alternatives: TFLearn, Lasagne, ...
- ▶ None of them are as **popular** as Keras!
- ▶ Will be **tightly integrated into TensorFlow** and officially supported by Google.
- ▶ Looks like a **safe future for Keras**!

 fchollet / keras


 Watch ▾ 954

 Unstar 13,448

 Fork 4,540





kli-nlpr commented on Jan 16

Contributor + 

Keras is gaining official Google support, and is moving into contrib, then core TF. If you want a high-level object-oriented TF API to use for the long term, Keras is the way to go.

<http://www.fast.ai/2017/01/03/keras/>

 1

 7

- ▶ Read the full story here: [Link](#)

Let's start!

- ▶ **How does the tutorial works?** You have the choice:
 1. You can just listen and learn from the code examples on the slides.
 2. You can follow along with the examples on your own laptop.
- ▶ **But** you'll learn most by taking the examples as starting point and play around at home.

Download all files:

```
git clone https://github.com/stwunsch/fermilab_keras_workshop
```

Set up the Python virtual environment:

```
cd fermilab_keras_workshop  
bash init_virtualenv.sh
```

Enable the Python virtual environment:

```
# This has to be done in every new shell!  
source py2_virtualenv/bin/activate
```

Keras Basics

Configure Keras Backend

- ▶ Two ways to configure Keras backend (Theano, TensorFlow or CNTK):
 1. Using **environment variables**
 2. Using **Keras config file** in `$HOME/.keras/keras.json`

Example setup using environment variables:

Terminal:

```
export KERAS_BACKEND=tensorflow
python your_script_using_keras.py
```

Inside a Python script:

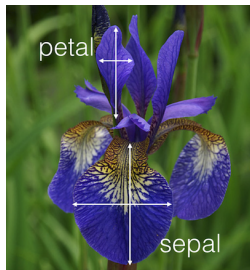
```
# Select TensorFlow as backend for Keras using enviroment variable `KERAS_BACKEND`
from os import environ
environ['KERAS_BACKEND'] = 'tensorflow'
```

Example Keras config using TensorFlow as backend:

```
$ cat $HOME/.keras/keras.json
{
    "image_dim_ordering": "th",
    "epsilon": 1e-07,
    "floatx": "float32",
    "backend": "tensorflow"
}
```

Example Using the Iris Dataset

- ▶ Next slides will introduce the basics of Keras using the example `fermilab_keras_workshop/keras/iris/train.py`.
- ▶ **Iris dataset:** Classify flowers based on their proportions
- ▶ **4 features:** Sepal length/width and petal length/width
- ▶ **3 targets** (flower types): Setosa, Versicolour and Virginica



Model Definition

- ▶ **Two types of models:** Sequential and the functional API
 - ▶ Sequential: Simply stacks all layers
 - ▶ Funktional API: You can do everything you want (multiple inputs, multiple outputs, ...).

```
# Define model
model = Sequential()

model.add(
    Dense(
        8, # Number of nodes
        kernel_initializer="glorot_normal", # Initialization
        activation="relu", # Activation
        input_dim=(4,) # Shape of inputs, only needed for the first layer!
    )
)

model.add(
    Dense(
        3, # Number of output nodes has to match number of targets
        kernel_initializer="glorot_uniform",
        activation="softmax" # Softmax enables an interpretation of the outputs as probabilities
    )
)
```

Model Summary

- ▶ `model.summary()` prints a description of the model
- ▶ **Extremely useful** to keep track of the number of free parameters

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 8)	40
dense_2 (Dense)	(None, 3)	27

Total params: 67
Trainable params: 67
Non-trainable params: 0

Define Loss Function, Optimizer, Validation Metrics...

- ▶ Everything is set in a single function, called the **compile** step of the model.
- ▶ Validation is performed after each training epoch (next slides).

```
# Set loss, optimizer and evaluation metrics
```

```
model.compile(  
    loss="categorical_crossentropy", # Loss function  
    optimizer=SGD(lr=0.10), # Optimizer algorithm  
    metrics=["accuracy",]) # Validation metric(s)
```

Data Preprocessing

- ▶ Some preprocessing steps are included in Keras, but mainly for text and image inputs.
- ▶ **Better option:** Using scikit-learn package ([Link to preprocessing module](#))
- ▶ **Single input** (4 features): [5.1, 3.5, 1.4, 0.2]
 - ▶ Needs to be scaled to the order of 1 to fit the activation function.
- ▶ **Single output** (3 classes): [1 0 0]
- ▶ **Common preprocessing:** Standardization of inputs
→ Operation: $\frac{\text{input} - \text{mean}}{\text{standard deviation}}$

Set up preprocessing

```
from sklearn.preprocessing import StandardScaler
preprocessing = StandardScaler()
preprocessing.fit(inputs)
inputs = preprocessing.transform(inputs)
```

Training

- ▶ Training is again a single call of the `model` object, called `fit`.

```
# Train  
model.fit(  
    inputs, # Preprocessed inputs  
    targets_onehot, # Targets in 'one hot' shape  
    batch_size=20, # Number of inputs used for  
                   # a single gradient step  
    epochs=10) # Number of cycles of the full  
               # dataset used for training
```

That's it for the training!

Training (2)

```
Epoch 1/10
150/150 [=====] - 0s 998us/step - loss: 1.1936 - acc: 0.2533
Epoch 2/10
150/150 [=====] - 0s 44us/step - loss: 0.9904 - acc: 0.5867
Epoch 3/10
150/150 [=====] - 0s 61us/step - loss: 0.8257 - acc: 0.7333
Epoch 4/10
150/150 [=====] - 0s 51us/step - loss: 0.6769 - acc: 0.8267
Epoch 5/10
150/150 [=====] - 0s 49us/step - loss: 0.5449 - acc: 0.8933
Epoch 6/10
150/150 [=====] - 0s 53us/step - loss: 0.4384 - acc: 0.9267
Epoch 7/10
150/150 [=====] - 0s 47us/step - loss: 0.3648 - acc: 0.9200
Epoch 8/10
150/150 [=====] - 0s 46us/step - loss: 0.3150 - acc: 0.9600
Epoch 9/10
150/150 [=====] - 0s 54us/step - loss: 0.2809 - acc: 0.9267
Epoch 10/10
150/150 [=====] - 0s 49us/step - loss: 0.2547 - acc: 0.9200
```

Save and Apply the Trained Model

Save model:

- ▶ Models are **saved as HDF5 files**: `model.save("model.h5")`
 - ▶ Combines description of weights and architecture in a single file
- ▶ **Alternative**: Store weights and architecture separately
 - ▶ Store weights: `model.save_weights("model_weights.h5")`
 - ▶ Store architecture: `json_dict = model.to_json()`

Load model:

```
from keras.models import load_model  
model = load_model("model.h5")
```

Apply model:

```
predictions = model.predict(inputs)
```

Wrap-Up

Training:

```
# Load iris dataset
# ...

# Model definition
model = Sequential()
model.add(Dense(8, kernel_initializer="glorot_normal", activation="relu", input_dim=(4,)))
model.add(Dense(3, kernel_initializer="glorot_uniform", activation="softmax"))

# Preprocessing
preprocessing = StandardScaler().fit(inputs)
inputs = preprocessing.transform(inputs)

# Training
model.fit(inputs, targets_onehot, batch_size=20, epochs=10)

# Save model
model.save("model.h5")
```

Application:

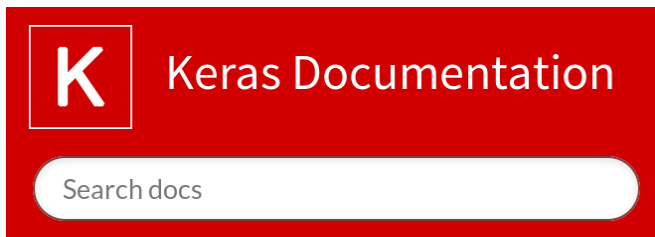
```
# Load model
model = load_model("model.h5")

# Application
predictions = model.predict(inputs)
```

That's a full training/application workflow in less than ten lines of code!

Available Layers, Losses, Optimizers, ...

- ▶ There's **everything you can imagine**, and it's **well documented**.
- ▶ Possible to **define own layers** and **custom metrics** in Python!
- ▶ Check out: `www.keras.io`



Advanced Usage of Keras

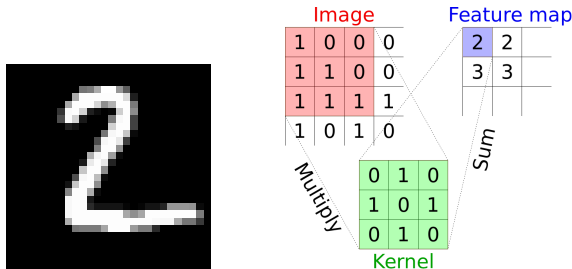
Example Using the MNIST Dataset

- ▶ Example in the repository:
`fermilab_keras_workshop/keras/mnist/train.py`
- ▶ **MNIST dataset?**
 - ▶ **Task:** Predict the number on an image of a handwritten digit
 - ▶ **Official website:** Yann LeCun's website ([Link](#))
 - ▶ Database of **70000 images of handwritten digits**
 - ▶ 28x28 pixels in greyscale as input, digit as label



- ▶ **Data format:**
 - ▶ **Inputs:** 28x28 matrix with floats in $[0, 1]$
 - ▶ **Target:** One-hot encoded digits, e.g., 2 \rightarrow $[0\ 0\ 1\ 0\ 0\ 0\ 0\ 0\ 0]$

Short Introduction to Convolutional Layers



- ▶ **Kernel:** Locally connected dense layer
- ▶ **Convolution:** Kernel moves similar to a sliding window over the image
- ▶ **Feature map:** Output “image” after application of the kernel

```
model = Sequential()

model.add(
    Conv2D(
        4, # Number of kernels/feature maps
        (3, # column size of sliding window used for convolution
         3), # row size of sliding window used for convolution
        activation="relu" # Rectified linear unit activation
    )
)
```

Model Definition

`fermilab_keras_workshop/keras/mnist/train.py:`

```
model = Sequential()

# First hidden layer
model.add(
    Conv2D(
        4, # Number of output filters or so-called feature maps
        (3, # column size of sliding window used for convolution
        3), # row size of sliding window used for convolution
        activation="relu", # Rectified linear unit activation
        input_shape=(28,28,1) # 28x28 image with 1 color channel
    )
)

# All other hidden layers
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Flatten())
model.add(Dense(16, activation="relu"))

# Output layer
model.add(Dense(10, activation="softmax"))

# Print model summary
model.summary()
```

Model Summary

- Detailed summary of model complexity with `model.summary()`

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 4)	40
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 4)	0
flatten_1 (Flatten)	(None, 676)	0
dense_1 (Dense)	(None, 16)	10832
dense_2 (Dense)	(None, 10)	170
Total params: 11,042		
Trainable params: 11,042		
Non-trainable params: 0		

Training With Validation Metrics

- ▶ Validation metrics are evaluated after each training epoch.
- ▶ In compile step, multiple **predefined validation metrics** can be booked, e.g., accuracy.
- ▶ **Custom metrics** are possible.

Booking a predefined metric:

```
# Compile model
model.compile(loss="categorical_crossentropy",
              optimizer=Adam(),
              metrics=["accuracy"])
```

Training with validation data:

```
model.fit(inputs, targets, validation_split=0.2) # Use 20% of the data for validation
```

```
Epoch 1/10
30000/30000 [=====] - 6s 215us/step - loss: 0.8095 - acc: 0.7565
- val_loss: 0.3180 - val_acc: 0.9085
Epoch 2/10
...
```

Training With Callbacks

- ▶ **Callbacks** are executed before and/or after each training epoch.
- ▶ Numerous **predefined** callbacks are available, **custom** callbacks can be implemented.

Definition of model-checkpoint callback:

```
# Callback for model checkpoints
checkpoint = ModelCheckpoint(
    filepath="mnist_example.h5", # Output similar to model.save("mnist_example.h5")
    save_best_only=True) # Save only model with smallest loss
```

Register callback:

```
model.fit(inputs, targets, # Training data
          batch_size=100, # Batch size
          epochs=10, # Number of training epochs
          callbacks=[checkpoint]) # Register callbacks
```

Training With Callbacks (2)

- ▶ Commonly used callbacks for improvement, debugging and validation of the training progress are implemented, e.g., **EarlyStopping**.
- ▶ Powerful tool: **TensorBoard** in combination with TensorFlow
- ▶ Custom callback: **LambdaCallback** or write callback class extending base class `keras.callbacks.Callback`

Callbacks

- Usage of callbacks
- Callback
- BaseLogger
- TerminateOnNaN
- ProgbarLogger
- History
- ModelCheckpoint
- EarlyStopping
- RemoteMonitor
- LearningRateScheduler
- TensorBoard
- ReduceLROnPlateau
- CSVLogger
- LambdaCallback
- Create a callback

Advanced Training Methods for Big Data

- ▶ The call `model.fit(inputs, targets, ...)` expects all inputs and targets to be already loaded in memory.
→ Physics applications have often data on Gigabyte to Terabyte scale!

These methods can be used to train on data that does not fit in memory.

- ▶ Training on **single batches**, performs a single gradient step:

```
model.train_on_batch(inputs, targets, ...)
```

- ▶ Training with data from a **Python generator**:

```
def generator_function():  
    while True:  
        yield custom_load_next_batch()
```

```
model.fit_generator(generator_function, ...)
```


Application on Handwritten Digits

- ▶ **PNG images of handwritten digits** are placed in `fermilab_keras_workshop/keras/mnist/example_images/`, have a look!



- ▶ Let's **apply our trained model** on the images:

```
./keras/mnist/apply.py keras/mnist/example_images/*.png
```

- ▶ **If you are bored on your way home:**
 1. Open with GIMP `your_own_digit.xcf`
 2. Dig out your most beautiful handwriting
 3. Save as PNG and run your model on it

Application on Handwritten Digits (2)

Predict labels for images:



```
keras/mnist/example_images/example_input_0.png : 7
keras/mnist/example_images/example_input_1.png : 2
keras/mnist/example_images/example_input_2.png : 1
keras/mnist/example_images/example_input_3.png : 0
keras/mnist/example_images/example_input_4.png : 4
keras/mnist/example_images/example_input_5.png : 1
keras/mnist/example_images/example_input_6.png : 4
keras/mnist/example_images/example_input_7.png : 9
keras/mnist/example_images/example_input_8.png : 6
keras/mnist/example_images/example_input_9.png : 9
```



Examples with Physics Data

Toy Calorimeter

- ▶ Data represent measurements in a toy-calorimeter
 - ▶ **Inputs:** 13 calorimeter layers with energy deposits
 - ▶ **Target:** Reconstruction of total energy deposit
- ▶ Example in repository:
`fermilab_keras_workshop/keras/calorimeter/train.py`



Implemented regression model:

```
model = Sequential()  
model.add(Dense(100, activation="tanh", input_dim=(13,)))  
model.add(Dense(1, activation="linear"))
```

- ▶ **Source:** [Link](#)

Deep Learning on the HIGGS Dataset

One of the most often cited papers about deep learning in combination with a physics application:

Searching for Exotic Particles in High-Energy Physics with Deep Learning

Pierre Baldi, Peter Sadowski, Daniel Whiteson

- ▶ **Topic:** Application of deep neural networks for separation of signal and background in an exotic Higgs scenario
- ▶ **Results:** Deep learning neural networks are more powerful than “shallow” neural networks with only a single hidden layer.

Let's reproduce this with minimal effort using Keras!

Deep Learning on the HIGGS Dataset (2)

Files:

- ▶ `fermilab_keras_workshop/keras/HIGGS/train.py`
- ▶ `fermilab_keras_workshop/keras/HIGGS/test.py`

Dataset:

- ▶ Number of events: 11M
- ▶ Number of features: 28

Shallow model:

```
model_shallow = Sequential()  
model_shallow.add(Dense(1000, activation="tanh", input_dim=(28,)))  
model_shallow.add(Dense(1, activation="sigmoid"))
```

Deep model:

```
model_deep = Sequential()  
model_deep.add(Dense(300, activation="relu", input_dim=(28,)))  
model_deep.add(Dense(300, activation="relu"))  
model_deep.add(Dense(300, activation="relu"))  
model_deep.add(Dense(300, activation="relu"))  
model_deep.add(Dense(300, activation="relu"))  
model_deep.add(Dense(1, activation="sigmoid"))
```

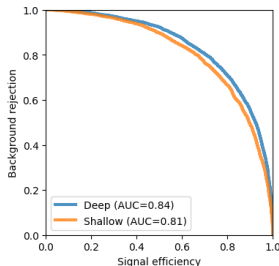
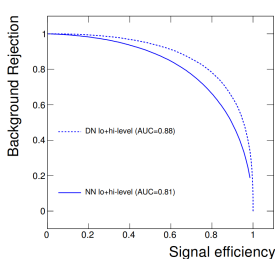
Training:

```
model.compile(loss="binary_crossentropy", optimizer=Adam(), metrics=["accuracy"])
```

```
model.fit(preprocessed_inputs, targets,  
        batch_size=100, epochs=10, validation_split=0.25)
```

Deep Learning on the HIGGS Dataset (3)

- Weights of deep and shallow model are part of the repository.



- Shallow model matches performance in the paper, but deep model can be improved.
→ **Try to improve it!** But you'll need a decent GPU...
- Keras allows to **reproduce this result with a total of 130 lines of code:**

```
# Count lines of code
$ wc -l keras/HIGGS/*.py
 62 keras/HIGGS/test.py
 68 keras/HIGGS/train.py
130 total
```

Useful Tools In Combination With Keras

TMVA Keras Interface

Prerequisites

- ▶ **Keras interface integrated in ROOT/TMVA since v6.08**
- ▶ Example for this tutorial is placed here:
`fermilab_keras_workshop/tmva/`
- ▶ You need ROOT with enabled PyROOT bindings. Easiest way to test the example is using CERN's **lxplus** machines:
 - ▶ `ssh -Y you@lxplus.cern.ch`
 - ▶ Source software stack LCG 91

How to source LCG 91 on lxplus:

```
source /cvmfs/sft.cern.ch/lcg/views/LCG_91/x86_64-slc6-gcc62-opt/setup.sh
```

Why do we want a Keras interface in TMVA?

1. **Fair comparison** with other methods
 - ▶ Same preprocessing
 - ▶ Same evaluation
2. **Try state-of-the-art DNN performance in existing analysis/application** that is already using TMVA
3. **Access data in ROOT files** easily
4. Integrate Keras in your **application** using **C++**
5. **Latest DNN algorithms in the ROOT framework** with **minimal effort**

How does the interface work?

1. **Model definition** done in **Python** using **Keras**
2. **Data management, training and evaluation** within the TMVA framework
3. **Application** using the TMVA reader



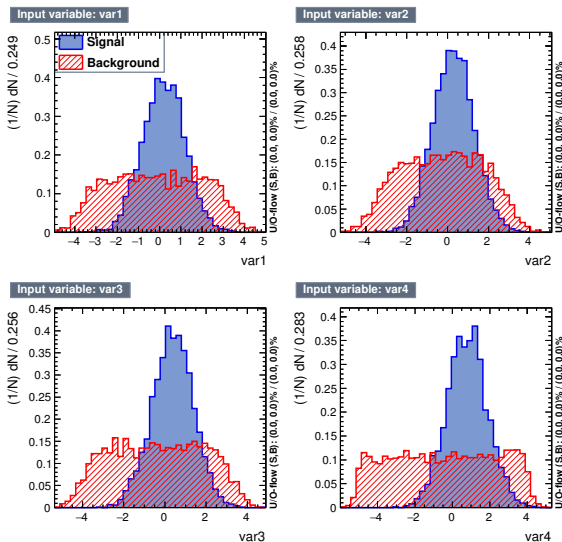
- The interface is implemented in the optional **PyMVA** part of TMVA:

```
# Enable PyMVA
```

```
ROOT.TMVA.PyMethodBase.PyInitialize()
```

Example Setup

- **Dataset** of this example is standard ROOT/TMVA test dataset for binary classification



Model Definition

- ▶ Setting up the model does not differ from using plain Keras:

```
model = Sequential()  
model.add(Dense(64, init='glorot_normal', activation='relu', input_dim=4))  
model.add(Dense(2, init='glorot_uniform', activation='softmax'))  
model.compile(loss='categorical_crossentropy', optimizer=Adam(), metrics=['accuracy',])  
model.save('model.h5')
```

- ▶ For **binary classification** the model needs **two output nodes**:

```
model.add(Dense(2, activation='softmax'))
```

- ▶ For **multi-class classification** the model needs **two or more output nodes**:

```
model.add(Dense(5, activation='softmax'))
```

- ▶ For **regression** the model needs a **single output node**:

```
model.add(Dense(1, activation='linear'))
```

Training

► Training options defined in the **TMVA booking options**:

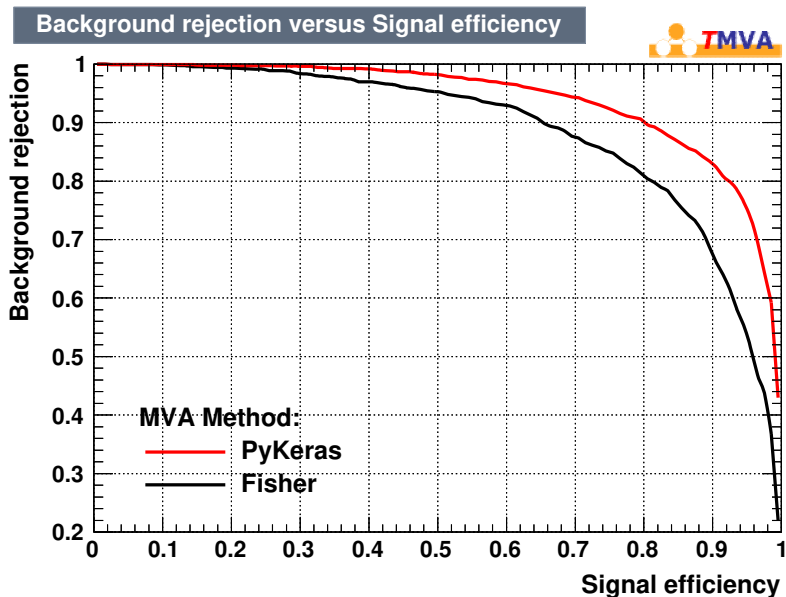
```
factory.BookMethod(dataloader, TMVA.Types.kPyKeras, 'PyKeras',  
    'H:V:VarTransform=G:' +  
    'Verbose=1'+\ # Training verbosity  
    'FilenameModel=model.h5:' +\ # Model from definition  
    'FilenameTrainedModel=modelTrained.h5:' +\ # Optional!  
    'NumEpochs=10:' +\  
    'BatchSize=32'+\  
    'ContinueTraining=false'+\ # Load trained model again  
    'SaveBestOnly=true'+\ # Callback: Model checkpoint  
    'TriesEarlyStopping=5'+\ # Callback: Early stopping  
    'LearningRateSchedule=[10,0.01; 20,0.001]')
```

That's it! You are ready to run!

```
python tmva/BinaryClassification.py
```

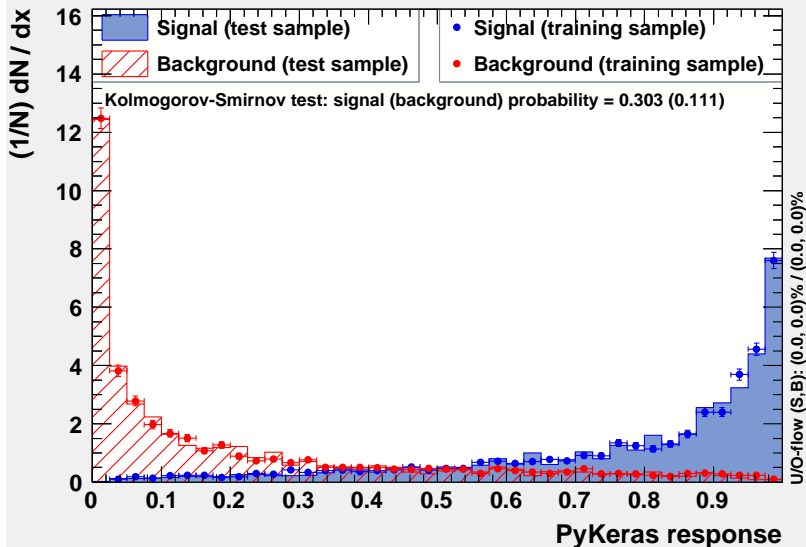
Run TMVA GUI to examine results: `root -l tmva/TMVAGui.C`

Training Results: ROC



Training Results: Overtraining Check

TMVA overtraining check for classifier: PyKeras



Application

- ▶ **Does not differ from any other TMVA method!**
- ▶ **Example** application can be found here:
fermilab_keras_workshop/
tmva/ApplicationBinaryClassification.py

Application (2)

Run `python tmva/ApplicationBinaryClassification.py`:

```
# Response of TMVA Reader
: Booking "PyKeras" of type "PyKeras" from
: BinaryClassificationKeras/weights/TMVAClassification_PyKeras.weights.xml.

Using Theano backend.
DataSetInfo      : [Default] : Added class "Signal"
DataSetInfo      : [Default] : Added class "Background"
                  : Booked classifier "PyKeras" of type: "PyKeras"
                  : Load model from file:
                  : BinaryClassificationKeras/weights/TrainedModel_PyKeras.h5

# Average response of MVA method on signal and background
Average response on signal:      0.78
Average response on background: 0.21
```

lwttn with Keras

What is lwttn?

- ▶ **Core problem:** TensorFlow and others are not made for event-by-event application!
- ▶ **C++ library** to apply neural networks
 - ▶ Minimal dependencies: C++11, Eigen
 - ▶ Robust
 - ▶ Fast
- ▶ **“Asymmetric” library:**
 - ▶ **Training** in any language and framework on any system, e.g., **Python and Keras**
 - ▶ **Application** in **C++** for real-time applications in a limited environment, e.g., high-level trigger
- ▶ **GitHub:** <https://github.com/lwttn/lwttn>
- ▶ **IML talk about lwttn by Daniel Guest:** [Link](#)
- ▶ **Tutorial** can be found here:
https://github.com/stwunsch/iml_keras_workshop

Load and Apply Model in C++ Using lwttn

Convert trained Keras model to lwttn JSON:

→ See the tutorial and README!

Load model:

```
// Read lwttn JSON config  
auto config = lwt::parse_json(std::ifstream("lwttn.json"));  
  
// Set up neural network model from config  
lwt::LightweightNeuralNetwork model(  
    config.inputs,  
    config.layers,  
    config.outputs);
```

Apply model:

```
// Load inputs from argv  
std::map<std::string, double> inputs;  
...  
  
// Apply model on inputs  
auto outputs = model.compute(inputs);
```